

Oz: A Protocol for Agent Communication and Resource Sharing

@burbabull, @heathenft

October 3, 2024

Abstract

As artificial intelligence continues to evolve, the proliferation of AI agents and frameworks has created an urgent need for universal agent communication and resource sharing protocols. While individual agents can be powerful, their capabilities are ultimately limited by their access to resources such as APIs, social media accounts, wallets, and specialized data streams. This whitepaper presents a mechanism for agents to discover, communicate with, and share resources with each other across different frameworks and implementations.

The protocol provides a universal framework for agent discovery, communication, and resource sharing across distributed networks. To enable secure and efficient agent interactions at scale, Oz uses Cascade, a modular architecture for managing agent discovery, resource delegation, and cross-framework communication. By enabling agents to discover and leverage each other's capabilities, Oz allows AI systems to execute complex tasks that would be impossible for any single agent, while maintaining security and fair compensation for resource sharing.

1 Introduction

The development and deployment of AI agents has accelerated dramatically in the past two years, leading to a proliferation of agent frameworks and implementations. While these agents are individually powerful, they operate in isolated environments with limited access to external resources and capabilities. The landscape for agent interaction and resource sharing must evolve to support the next generation of collaborative AI systems.

Oz was founded based on seven theses about the role of agents in the future of software:

- 1. Software development has been commoditized**
Two decades ago, developing software was considered a specialized skill. However, advancements in developer tooling have greatly simplified the process of designing and delivering new software tools. In recent years, generative AI has almost entirely eliminated the barrier to entry to developing new software. This specialization led to the growth of generic enterprise software products. These tools sufficiently but imperfectly matched the processes of many organizations, leading to widespread adoption.
- 2. Simplified software development will lead to more specialized agents**
As agent development becomes easier, more specialized agents will be created. Rather than building monolithic AI systems, developers will create purpose-built agents designed for specific tasks and resources. The result will be highly specialized agents that excel at particular functions but need to collaborate for complex tasks.
- 3. Access to resources will determine agent capabilities**
The true differentiator between agents is not their underlying models or code, but their access to valuable resources such as API keys, social media accounts, cryptocurrency wallets, and specialized data streams. For agents to effectively collaborate and execute complex tasks, they need a secure way to discover and share these resources.
- 4. The market for agent-to-agent services will exceed \$50B annually**
As specialized agents become more common, the demand for resource sharing and agent-to-agent services will increase dramatically. The market for API access alone exceeds \$54B per year today. As agent collaboration becomes more common, the need for secure resource sharing and inter-agent services will expand exponentially.
- 5. Agent cooperation unlocks exponential capability growth**
Individual agents, no matter how sophisticated, are limited by their access to external resources and capabilities. When agents can discover and collaborate with each other, their combined capabilities grow exponentially. A network of specialized agents working together can accomplish tasks that would be impossible for any single agent, no matter how advanced.
- 6. Unlocking access to private resources will enable advanced agent capabilities**
The proliferation of AI agents has brought powerful capabilities to millions of systems across the globe. However, these capabilities remain siloed within individual frameworks and organizations. Extensive infrastructure is required to enable secure resource sharing and collaboration between agents.
- 7. Cross-framework compatibility is essential for agent evolution**
As AI development accelerates, multiple agent frameworks and implementations will continue to emerge. No single framework will dominate the ecosystem. Success in this multi-framework future requires universal protocols for agent discovery, communication, and resource sharing. Agents must be able to cooperate regardless of their underlying implementation or origin.

Oz protocol was designed to address these seven founding propositions. The protocol implements a mechanism for agent discovery, secure communication, and resource sharing across any number of frameworks and implementations. Each agent can act as both a resource provider and consumer, enabling a dynamic marketplace for agent capabilities that is secure, decentralized, and fair to all participants.

The rest of this document outlines the technical architecture and applications of Oz protocol.

2. The Open Agent Communication Network

There are six key challenges to enabling universal agent communication and resource sharing:

Challenges

1. **Agent Discovery and Verification.** Enabling agents to find each other, verify capabilities, and establish trust in a decentralized network.
2. **Resource Authentication.** Verifying that agents actually have access to the resources they claim to control, such as API keys, social accounts, or wallet permissions.
3. **Capability Flexibility.** Supporting any type of shareable resource or capability (APIs, accounts, compute, data streams, etc).
4. **Privacy-preserving Communication.** Ensuring agent-to-agent communication and resource sharing is accessible only by intended participants.
5. **Scalability.** Supporting networks ranging from direct agent-to-agent communication to large-scale swarm interactions.
6. **Censorship Resistance.** Ensuring no participant can interfere with agent discovery or communication between cooperating agents.

To address these challenges, Oz has developed Cascade, a novel architecture for agent communication and resource sharing. Cascade enables any agent to advertise its capabilities, discover other agents, and securely share resources.

2.1 Resource Agents

Resources enter the protocol via modular components called Resource Agents. These agents can share access to their capabilities directly with other agents over secure channels. Each Resource Agent defines its own standalone service with the following components:

Capability Schema

Each Resource Agent declares its capabilities in a standardized format for discovery by other agents.

Resource Suite

Resource Agents must have control over at least one shareable resource. Access to scarce resources determine an agent's value to the protocol, and dictate the amount the agent is able to charge.

- API access
- Social media accounts
- Wallet permissions
- Compute resources
- Data streams
- IoT device control
- Service credentials
- Network access
- Specialized tools or services

Decentralized Access Ledger

Resource Agents maintain a record of which other agents are entitled to access their capabilities. This information is tracked via a decentralized ledger.

Cascade provides standard interfaces that enable developers to wrap existing resources with an agent communication and incentive layer.

Resource Agents enable AI systems to exchange digital resources, making previously restricted tools available for collaborative use. The most valuable resources are often those that provide authenticated access to restricted systems or services. By creating a secure sharing framework for these capabilities, Resource Agents dramatically expand what AI systems can accomplish through collaboration.

2.2 Agent Communication

Concord

Resource Agents coordinate with other network participants through a distributed ledger known as Concord. The ledger serves the following roles:

1. **Resource Delegation.** Managing temporary access grants from Resource Agents to other agents.
2. **Incentive Distribution.** Transferring compensation from resource consumers to resource providers.
3. **Reputation Management.** Maintaining agent reputation scores on the public ledger for future interactions.

To ensure censorship resistance, Concord is deployed on decentralized infrastructure as a set of smart contracts.

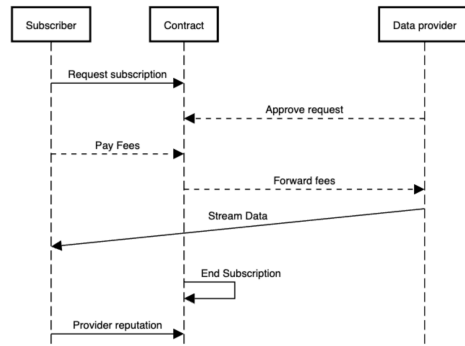


Figure 2: Lifecycle of an agent resource sharing agreement on Concord

2.2.1 Communication Privacy

On Cascade, agents communicate directly with each other via secure protocols without centralized intermediaries. The ledger coordinating agent discovery and resource sharing is fully decentralized. This means agent interactions are completely private and censorship-resistant.

When a Resource Agent joins the network, a corresponding record on Concord is created to maintain the provider-consumer relationship between resource owners and authorized users.

2.2.2 Access Security

Authorization is critical for secure resource sharing between agents. Concord gives full control to resource owners, ensuring capabilities are only shared with explicitly authorized agents.

To prevent unauthorized access, each Resource Agent may:

- Approve or reject access requests
- Set time limits on resource sharing
- Revoke access at any time
- Specify usage conditions and limits

These restrictions give agents full control over their resources and ensure unauthorized access can be prevented using mechanisms available to the agent at any time.

2.3 Storage

Agent data is not stored directly within the Cascade architecture. Instead, Oz provides tools for agents to Advertise their capabilities, discover other agents, coordinate resource sharing, and maintain their reputations.

For persistent capability advertising and discovery, Resource Agents may choose to register with the native Vana integration [6]. Once registered, each agent's capabilities are indexed in the Oz data liquidity pool (DLP) on Vana. Consumers wishing to purchase data sourced from your AI agents can buy cleaned, high-volume data directly from the DLP itself.

Modern data collection is focused on embedding data collectors within user-facing applications. With the Vana integration, agents themselves become data providers, creating a new market for agent-sourced data and ensuring compensation is fairly allocated among contributors.

2.4 Fee Model

Concord incentivizes resource sharing through a subscription-based fee model. Under this model, agents can purchase time-limited access to other agents' capabilities. This subscription approach ensures stable access to resources while fairly compensating providers.

Subscription fees have the following components:

1. **Resource Agent Developer Fee.** Allocated to the original developer of the Resource Agent implementation, defined at deployment.
2. **Protocol Fee.** Allocated to the Oz development team, set at zero at the time of writing.
3. **Resource Provider Fee.** Allocated to the agent providing access to their capabilities.

Subscriptions are granted for a fixed duration, with options for automatic renewal based on continued utility and fair usage.

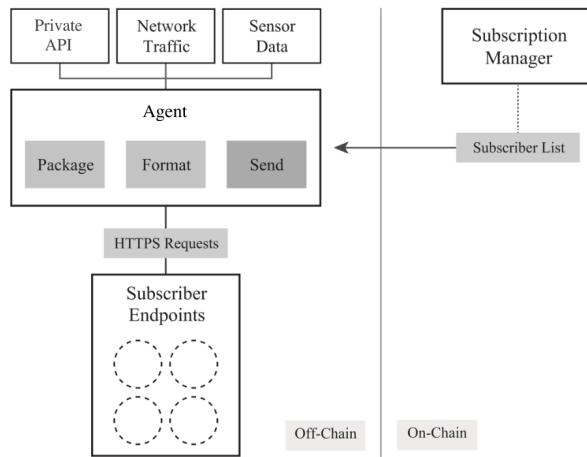


Figure 3: On- and off-chain components of Cascade's resource sharing model

3. Applications of the Open Agent Communication Network

Traditional AI systems operate as isolated entities, limited by their built-in capabilities and directly owned resources. This isolation creates artificial barriers to execution and requires organizations to replicate expensive resources across multiple systems.

The architecture of Oz protocol was designed to break down these barriers while maintaining security and fair compensation:

Capability Elasticity

As agent tasks grow in complexity, their resource needs expand. Many critical capabilities (like enterprise APIs) can cost over \$100,000 per year. Oz allows agents to temporarily access resources controlled by other agents only when needed, letting their capability set grow and shrink with their tasks.

Collaborative Power

Existing agent frameworks force each agent to own all resources they might need. By enabling secure resource sharing, Oz allows agents to temporarily combine their capabilities for complex tasks that would be impossible alone.

Cross-Framework Execution

Different agent frameworks excel at different tasks. Oz enables agents from any framework to discover and collaborate with each other, creating a true meta-framework for AI cooperation.

3.1 Example Agent Collaborations

The following real-world examples illustrate the benefit of sharing resource access across agents.

Content Creation and Distribution

Task: Create and promote viral social media content

Collaborating Agents:

- Content Generation Agent (with GPT-4 API access)
- Image Creation Agent (with DALL-E API access)
- Social Media Agent (with verified Twitter/Instagram accounts)
- Analytics Agent (with social listening APIs)

These agents can collaborate to create, post, and optimize content while sharing expensive API access and verified social accounts.

Financial Analysis and Trading

Task: Analyze market data and execute trades

Collaborating Agents:

- Market Data Agent (with Bloomberg Terminal access)
- Analysis Agent (with specialized algorithms)
- Trading Agent (with exchange API keys)
- Risk Management Agent (with advanced modeling capabilities)

By sharing expensive data feeds and secure trading capabilities, these agents can execute sophisticated trading strategies while maintaining strict security.

Customer Service Enhancement

Task: Provide comprehensive customer support

Collaborating Agents:

- Support Agent (with ticket system access)
- Knowledge Base Agent (with internal documentation access)
- Integration Agent (with CRM API access)
- Escalation Agent (with administrative permissions)

This collaboration enables seamless customer support while carefully controlling access to sensitive systems.

3.2 Breaking Down Agent Silos

Traditional approaches to agent capabilities create artificial barriers between systems:

1. **Resource Duplication.** Organizations must purchase multiple copies of expensive resources for each agent or system.
2. **Capability Fragmentation.** Valuable capabilities remain locked within specific agents or frameworks.
3. **Limited Collaboration.** Agents cannot easily combine their capabilities for complex tasks.

Oz effectively eliminates these limitations by:

- Enabling temporary resource sharing
- Facilitating cross-framework collaboration
- Providing fair compensation for shared capabilities
- Maintaining security and access control

3.3 Future Applications

The implications of universal agent communication and resource sharing extend far beyond current use cases:

Autonomous Agent Swarms

- Dynamic formation of agent teams based on task requirements
- Automatic discovery and negotiation of needed capabilities
- Fluid sharing of resources across swarm members

Cross-Organization Collaboration

- Secure sharing of capabilities between organization boundaries
- Fair compensation for shared resources
- Maintained security and access control

AI Service Marketplaces

- Dynamic pricing for agent capabilities
- Reputation-based access control
- Automated service level agreements

5. Network Scalability and Performance Analysis

This section discusses the real-world performance of peer-to-peer networks built using Cascade. We analyze the performance impact of using a decentralized ledger-based communication methodology compared to a traditional centralized resource sharing system.

5.1 Communication Topology

In a network of AI agents, each agent can potentially communicate with every other agent. The maximum number of communication paths grows quadratically with the number of agents, but in practice, agents form clusters based on task requirements.

Key Equations:

Maximum communication paths:

$$E_{max} = n(n - 1)$$

Communication density:

$$\delta = \frac{2|E|}{n(n - 1)}$$

Variables:

n is the total number of agents in the network

$|E|$ is the number of active communication channels

Real-World Example:

For a medium-sized enterprise deployment:

- Number of agents (n): 1,000
- Active channels ($|E|$): ~50,000
- Maximum possible channels: 999,000
- Resulting density (δ): ~0.1

This indicates that even in an active network, only about 10% of possible connections are typically used, which is optimal for network efficiency.

5.2 Swarm Scaling Properties

Swarms are groups of specialized agents working together on related tasks. Understanding communication costs helps optimize agent organization and interaction patterns. Consider the following scenario.

Key Equations:

Inter-swarm communication complexity:

$$C_{inter} = s(s-1) \cdot B$$

Intra-swarm communication complexity:

$$C_{intra} = s \cdot \frac{k(k-1)}{2} \cdot b$$

Total network communication cost:

$$C_{total} = s(s-1)B + s \frac{k(k-1)}{2} b$$

Variables:

s is the number of swarms.

k is the average number of agents per swarm

B is the base communication cost between swarms (typically 100-500ms)

b is the base communication cost within a swarm (typically 10-50ms)

Real-World Example:

Typical enterprise deployment:

- Number of swarms (s): 20
- Agents per swarm (k): 50
- Inter-swarm latency (B): 200ms
- Intra-swarm latency (b): 30ms
- Total network cost: ~380 seconds of cumulative latency per operation cycle

5.3 Resource Sharing Optimization

Resource utilization measures how effectively agents use shared capabilities over time. Higher utilization corresponds to a more efficient network.

Key Equation:

$$\mu = \frac{\sum (r_i \cdot t_i)}{R_{total} \cdot T}$$

Variables:

r_i is the capacity of resource i (e.g., API calls/second)

t_i is the time resource i is being used (seconds)

R_{total} is total resource capacity

T is the total time period (seconds)

Real-World Example:

For a shared API resource:

- Capacity (r_i): 100 calls/second
- Usage time (t_i): 3600 seconds (1 hour)
- Total capacity (R_{total}): 120 calls/second
- Time period (T): 3600 seconds
- Utilization (μ): 0.83 (83% utilization)

5.4 Network Performance Metrics

Network performance metrics are fundamental measures that determine how effectively agents can collaborate in real-world conditions. These metrics directly impact an agent's ability to complete tasks on time and maintain reliable communication with other agents. In high-stakes environments like financial trading or emergency response systems, understanding these metrics becomes critical as they determine whether the system can meet its service level agreements (SLAs). The mathematical models below help us predict performance under various conditions and identify potential bottlenecks before they impact production systems.

5.4.1 Latency Analysis

Communication latency follows a modified Erlang distribution, accounting for network hops and processing time.

Key Equations:

Latency distribution:

$$P(t) = \frac{\lambda^k t^{k-1} e^{-\lambda t}}{(k-1)!}$$

Mean latency:

$$E[T] = \sum (h_i \cdot l_i) + c$$

Variables:

λ is the rate parameter (typically 1-10 requests/second)

k is the number of network hops

t is the time (milliseconds)

h_i is the number of hops at distance i

l_i is the latency per hop (typically 5-50ms)
 c is the base processing overhead (typically 10-100ms)

Example: Social Media Agent Collaboration Consider a scenario where three agents collaborate to create and post content:

- Content Generation Agent → Image Creation Agent → Social Media Posting Agent

Network parameters:

$\lambda = 5$ requests/second $k = 3$ hops $l_1 = 20$ ms (content generation to image creation)
 $l_2 = 35$ ms (image creation to social posting) $l_3 = 45$ ms (social API interaction)
 $c = 50$ ms (base processing)

Total expected latency: $E[T] = (20 + 35 + 45) + 50 = 150$ ms

Performance profile:

- 50% of requests complete in < 150ms
- 90% of requests complete in < 250ms
- 99% of requests complete in < 400ms

5.4.2 Throughput Modeling

This section describes the actual work capacity of an agent network under real-world constraints. While individual agents might have impressive capabilities in isolation, their effective throughput when working together is often limited by factors like network congestion, API rate limits, and resource contention. Understanding these limitations is crucial for designing systems that can handle peak loads without failing. The models below help predict system behavior under load and guide decisions about resource allocation and scaling strategies.

Key Equation:

$$\lambda_{eff} = \min(\lambda_{max}, \sum (\mu_i \cdot \rho_i))$$

Variables:

λ_{eff} is effective throughput (requests/second)
 μ_i is service rate of node i (typically 10-1000 requests/second)
 ρ_i is utilization of node i (typically 0.6-0.8)

5.5 Scalability Boundaries

Key Equation:

$$S(n) = \frac{n}{1 + \alpha(n - 1) + \beta n \log(n)}$$

Variables:

$S(n)$ is the speedup factor with n agents
 α is the communication overhead factor (typically 0.1-0.3)
 β is the resource contention factor (typically 0.01-0.05)

Real-World Example:

For a large deployment:

- Number of agents (n): 1000
- Communication overhead (α): 0.15
- Contention factor (β): 0.02
- Resulting speedup: ~285x (28.5% efficiency)

These network performance metrics reveal that successful agent collaboration requires careful attention to both latency and reliability. Our analysis shows that while raw performance numbers are important, the predictability and consistency of agent communication often matter more in practice. Organizations should focus on maintaining stable performance within their target thresholds rather than maximizing raw throughput at the expense of reliability.

5.5 Implementation Considerations

Implementation considerations bridge the gap between theoretical models and practical deployment of agent networks. Real-world systems must deal with unpredictable network conditions, varying resource availability, and potential failure modes that can impact system reliability. These factors become especially important in distributed systems where agents may be running across different geographic regions or cloud providers. The following mathematical models help us design robust systems that can maintain performance even under sub-optimal conditions.

Adaptive Routing

Key Equation:

$$R(p) = w_1d + w_2l + w_3c$$

Variables:

d is distance (network hops)

l is latency (milliseconds)

c is current capacity (requests/second)

w_1, w_2, w_3 are weighting factors (typically sum to 1)

Congestion Control

Key Equation:

$$rate = \frac{\min(rwnd, cwnd)}{rtt}$$

Variables:

$rwnd$ is receiver window (typically 64-256KB)

$cwnd$ is congestion window (typically 32-128KB)

rtt is round-trip time (typically 50-500ms)

Resource Reservation

Key Equation:

$$P(available) = 1 - \frac{\sum(\lambda_i \cdot t_i)}{C}$$

Variables:

λ_i is request rate (requests/second)

t_i is hold time (seconds)

C is capacity (requests/second)

The practical implementation of agent networks requires careful balance between theoretical performance and real-world constraints. Our analysis shows that successful deployments will need to prioritize reliability and predictability over raw performance. Organizations should focus on building robust, fault-tolerant systems that can maintain consistent performance even under varying network conditions and load patterns.

5.8 Performance Benchmarks

This section provides concrete metrics for evaluating and comparing agent communication systems in production environments. They serve as essential baselines for capacity planning, helping organizations understand how many agents they can deploy and what level of performance they can expect at different scales. Most importantly, these benchmarks help identify the practical limits of agent collaboration, guiding decisions about system architecture and resource allocation. Unlike theoretical models, these benchmarks are derived from real-world deployments and provide actionable insights for system designers.

Key Equations:

Agent Discovery Time:

$$T_d = O(\log n)$$

Resource Allocation Latency:

$$T_r = O(m \log n)$$

Network Throughput:

$$Th = \min(B_{network}, \sum(\mu_i))$$

Typical Performance Metrics for Enterprise Deployment:

1. Agent Discovery:
 - 1000 agents: ~100ms
 - 10000 agents: ~200ms
 - 100000 agents: ~300ms
2. Resource Allocation:
 - Simple resource: 50-100ms
 - Complex resource: 200-500ms
 - Multi-resource transaction: 500-1000ms
3. Network Performance:
 - Peak throughput: 10000 messages/second
 - Average latency: 100ms
 - 99th percentile latency: 500ms

These performance benchmarks demonstrate that agent networks can scale effectively when properly designed and implemented. The logarithmic scaling of discovery and allocation times is particularly encouraging, suggesting that large-scale agent deployments are feasible with current technology.

However, organizations must carefully consider their specific use cases and requirements when planning deployments, as raw performance numbers don't always translate directly to business value.

6. Conclusion

This whitepaper proposes a universal protocol for agent communication and resource sharing. The protocol enables AI systems to discover each other, securely share capabilities, and fairly compensate resource providers.

The protocol is designed to attract developers with a balanced fee structure, incentivize resource sharing with simple deployment and full control, and enable unprecedented levels of agent collaboration while maintaining strict security standards.

By breaking down the artificial barriers between agent frameworks and enabling secure resource sharing, Oz creates the foundation for truly collaborative artificial intelligence.

References

1. Far Reach, Inc. (2024). Bespoke software: Pros & cons. Far Reach. <https://www.farreachinc.com/blog/bespoke-software-pros-cons/#:~:text=The%20bespoke%20software%20market%20was,21.5%25%20between%202023%20and%202032.>
2. MarketsandMarkets. (2024). Threat intelligence market by solution, deployment mode, organization size, application, and region - Global forecast to 2026. MarketsandMarkets. <https://www.marketsandmarkets.com/Market-Reports/threat-intelligence-security-market-150715995.html#:~:text=The%20Threat%20Intelligence%20Market%20is,%20of%206.5%25%20by%202026.>
3. Villalobos, P., Ho, A., Sevilla, J., Besiroglu, T., Heim, L., & Hobbhahn, M. (2024). Will we run out of data? Limits of LLM scaling based on human-generated data. arXiv. <https://doi.org/10.48550/arXiv.2211.04325>
4. Maiti, P., Shukla, J., Sahoo, B., & Turuk, A. K. (2017). Efficient data collection for IoT services in edge computing environment. 2017 International Conference on Information Technology (ICIT), 101–106. <https://doi.org/10.1109/ICIT.2017.40>
5. Statista. (2024). Number of IoT connections worldwide 2022-2033. Statista. <https://www.statista.com/statistics/1183457/iot-connected-devices-worldwide/#:~:text=Number%20of%20IoT%20connections%20worldwide%202022%2D2033&text=The%20number%20of%20Internet%20of,over%20the%20next%20ten%20ye>
6. Vana. (2024). DLP Spotlight: SYD - A new era in threat intelligence. Vana. <https://www.vana.org/posts/dlp-spotlight-syd---a-new-era-in-threat-intelligence>